

# Project Master Summary: Bitcoin Sentiment & Trader Performance

Created by Nishant Sahni

September 25, 2025

## Contents

<b>1 Project Objective</b>	<b>3</b>
<b>2 Raw Inputs</b>	<b>3</b>
<b>3 Produced Files Inventory</b>	<b>3</b>
<b>4 Data Cleaning Steps</b>	<b>4</b>
4.1 Trades Cleaning . . . . .	4
4.2 Sentiment Cleaning . . . . .	5
4.3 Join Criteria . . . . .	5
<b>5 Reproducible Run Commands</b>	<b>5</b>
<b>6 Key Numeric Descriptive Results</b>	<b>5</b>
6.1 Dataset Scale . . . . .	5
6.2 Closed PnL (trade-level) Stats . . . . .	6
<b>7 Event Study</b>	<b>6</b>
7.1 Methodology . . . . .	6
7.2 Counts . . . . .	6
7.3 Mean Daily PnL at t=0 . . . . .	6
<b>8 Granger-Causality Analysis</b>	<b>7</b>
8.1 Preprocessing . . . . .	7
8.2 Market-Level Results (p-values for F-test by lag) . . . . .	7
8.3 Interpretation . . . . .	7
<b>9 Notebooks &amp; Code Structure</b>	<b>7</b>
<b>10 Final Detailed Report Contents (final_report_detailed.pdf)</b>	<b>7</b>
<b>11 Limitations, Biases, and Validity Threats</b>	<b>8</b>
<b>12 Recommended Next Analyses</b>	<b>8</b>
12.1 Panel Fixed-Effects Regression . . . . .	8
12.2 Predictive Model (Classification) . . . . .	9
12.3 Trader Clustering (Segmentation) . . . . .	9
12.4 Event Study Variants . . . . .	9

<b>13</b>	<b>How to Incorporate BTC Price &amp; Volatility</b>	<b>9</b>
<b>14</b>	<b>Full Prioritized Action Plan</b>	<b>9</b>
<b>15</b>	<b>Concrete Recommendations for a Trading Team</b>	<b>10</b>
<b>16</b>	<b>Exact Code Snippets &amp; Helper Utilities</b>	<b>10</b>
16.1	Auto-detect DATA_ROOT . . . . .	10
16.2	Load Cleaned Data . . . . .	10
16.3	Event Study Helper . . . . .	11
16.4	Granger Causality Quick Call . . . . .	11
<b>17</b>	<b>Location of Important Outputs</b>	<b>11</b>

# 1 Project Objective

Explore the relationship between Bitcoin market sentiment (Fear / Greed index) and Hyperliquid traders' realized performance — find patterns, check predictability, quantify effects, and produce reproducible deliverables (cleaned data, notebooks, analysis, event studies, regressions, report, presentation, and a demo).

## 2 Raw Inputs

- `/mnt/data/historical_data.csv`: Hyperliquid trades (per-trade records).
  - **Key raw columns detected**: Account, Coin, Execution Price, Size Tokens, Size USD, Side, Timestamp IST, Timestamp, Closed PnL, Trade ID, Order ID, Transaction Hash, Start Position, Direction, Fee, Crossed.
- `/mnt/data/fear_greed_index.csv`: Bitcoin Fear/Greed index.
  - **Columns detected**: timestamp, value (0–100 numeric index), classification (Extreme Fear/Fear/Neutral/Greed/Extreme Greed), date.

## 3 Produced Files Inventory

All generated files are saved under: `/mnt/data/submission_package/`.

- **Data (cleaned)**
  - `trades_clean.csv`: Cleaned & normalized trades (UTC time, date, size\_abs, closedPnL, is\_win, side\_std, closedPnL\_outlier, etc.)
  - `sentiment_clean.csv`: Cleaned sentiment (date, value, sent\_index\_raw, sent\_index\_norm [-1..1], sent\_label\_clean, sent\_numeric\_label)
- **Aggregates / Features**
  - `daily_account_agg.csv`: Per-account, per-day aggregates (daily\_pnl, num\_trades, win\_rate, avg\_size, median\_size, pnl\_std, num\_outliers)
  - `daily_overall_agg.csv`: Market-level daily aggregates (sum/averages)
  - `merged_daily_sentiment.csv`: Merged daily market metrics with sentiment
- **Event Study Outputs**
  - `event_study_market_low.csv` & `event_study_market_high.csv`: Market-level event stats (rel\_day -3..3)
  - `event_study_account_low.csv` & `event_study_account_high.csv`: Account-level event stats
  - `event_study_market_low.png`, `event_study_market_high.png`, `event_study_account_low.png`, `event_study_account_high.png`: PNG plots
- **Causality**
  - `granger_full_summary.json`: ADF results, diff flags, Granger p-values serialized

- granger\_account\_summary.csv: Per-account Granger summary (top tested accounts, p-values)
- **Notebooks & Code**
  - 1\_eda.ipynb
  - 2\_features\_and\_aggregation.ipynb
  - 3\_analysis.ipynb
  - streamlit\_app.py: Tiny demo to inspect tables
- **Deliverables**
  - final\_report.pdf: Short 2-page report
  - final\_report\_detailed.pdf: Detailed multi-page final report
  - presentation.pdf: 8-slide deck
  - README.md
  - requirements.txt

## 4 Data Cleaning Steps

### 4.1 Trades Cleaning

- **Column normalization:** Renamed columns to lowercase/standard names (e.g., Closed PnL -> closedPnL).
- **Timestamp parsing:**
  - Preferred timestamp\_ist (assumed IST), localized to Asia/Kolkata, then converted to **UTC**.
  - Fell back to timestamp (parsed as UTC) if timestamp\_ist was unavailable.
  - Final column: time (timezone-aware UTC). Derived date from time.
- **Numeric coercions:** Coerced execution\_price, size\_tokens, size\_usd, closedPnL, fee to numeric, setting errors to NaN.
- **Derived size\_abs:** Preferred size\_tokens.abs(); fell back to size\_usd.abs().
- **Standardize side:** Mapped common values to side\_std (**+1** for buy/long, **-1** for sell/short). Inferred from size sign if contradictory.
- **Flags & derived fields:**
  - is\_win = closedPnL > 0
  - closedPnL\_abs, closedPnL\_outlier flag where absolute PnL > 99.9th percentile.
- **Row removals:**
  - Dropped rows missing time, account, or size\_abs.
  - Removed zero-size trades.
  - Removed duplicates by trade\_id (kept last).

## 4.2 Sentiment Cleaning

- **Timestamp parsing:** Final date field is the calendar date.
- **Value normalization:**
  - `sent_index_raw`: Retained original 0–100 value.
  - `sent_index_norm`: Normalized to [-1, 1] using the formula:

$$\text{sent\_index\_norm} = \frac{(\text{sent\_index\_raw} - \min)}{(\max - \min)} \times 2 - 1$$

- **Labels:** Kept classification as `sent_label_clean` and mapped to numeric labels (`sent_numeric_label`): Extreme Fear: -2, Fear: -1, Neutral: 0, Greed: 1, Extreme Greed: 2.
- **Deduplication:** Kept the last entry per date and forward-filled numeric values where appropriate.

## 4.3 Join Criteria

Daily aggregations were merged on date (UTC date) using a left join.

## 5 Reproducible Run Commands

Run from the parent directory of `submission_package`.

```
1 # from parent folder of submission_package
2 python -m venv venv
3 source venv/bin/activate
4 pip install -r submission_package/requirements.txt
5 # (optional) pillow for embedding/PNG handling:
6 pip install pillow
7
8 # Launch Jupyter to run notebooks:
9 jupyter notebook
10 # open and run:
11 # submission_package/1_eda.ipynb
12 # submission_package/2_features_and_aggregation.ipynb
13 # submission_package/3_analysis.ipynb
14
15 # Launch demo:
16 streamlit run submission_package/streamlit_app.py
```

## 6 Key Numeric Descriptive Results

### 6.1 Dataset Scale

- **Total trades processed:** 211,224
- **Total sentiment rows:** 2,644
- **Time span (trades):** 2023-01-04 19:36:00 UTC → 2025-12-04 18:25:00 UTC
- **Critical data-quality note:** The max trade timestamp (2025-12-04) is in the future relative to the analysis date (2025-09-25) and must be verified.

- **Unique accounts:** 32
- **Unique symbols:** 246

## 6.2 Closed PnL (trade-level) Stats

- **Sum of closedPnL:** \$10,296,958.94
- **Mean closedPnL:** \$48.7490
- **Median closedPnL:** \$0.00
- **Standard Deviation:** \$919.165
- **Skewness:** 30.6994 (extremely heavy positive skew)
- **Interpretation:** A tiny fraction of trades generates most of the realized PnL. The median of 0 indicates many trades are break-even. Analysis must use robust statistics (median, trimmed mean) or winsorize.

## 7 Event Study

### 7.1 Methodology

- **Event Definition:** Extreme events were defined as the bottom 10% (Extreme Fear) and top 10% (Extreme Greed) of the `sent_index_norm`.
- **Window:** For each event date  $t$ , metrics were computed for the window  $t-3, \dots, t+3$ .
- **Aggregation Lenses:**
  1. **Market-level:** Sum of `daily_pnl` across all accounts.
  2. **Account-level:** Average of each account's `daily_pnl` for that day, then averaged across all events.
- **Statistics:** For each relative day, the mean, standard deviation, sample count, and 95% CI were computed.

### 7.2 Counts

- **Extreme Fear days** (bottom 10%): 48 unique event dates
- **Extreme Greed days** (top 10%): 58 unique event dates

### 7.3 Mean Daily PnL at $t=0$

- **Market-level (Extreme Fear):** \$26,418.35
- **Market-level (Extreme Greed):** \$31,416.26
- **Account-level (Extreme Fear):** \$2,929.08
- **Account-level (Extreme Greed):** \$3,098.98
- **Interpretation:** On average, event days (both Fear and Greed) coincide with elevated PnL, suggesting increased activity or intensity. However, wide confidence intervals indicate large dispersion.

## 8 Granger-Causality Analysis

### 8.1 Preprocessing

- Daily series used: `market_daily_pnl` and `sent_index_norm`.
- **Stationarity:** Augmented Dickey-Fuller (ADF) tests were performed. Non-stationary series were differenced once.
- **Test:** `statsmodels.tsa.stattools.grangercausalitytests` was used with lags from 1 to 7.

### 8.2 Market-Level Results (p-values for F-test by lag)

- **Sentiment** → **market\_daily\_pnl**:
  - lag1: 0.1251, lag2: 0.1930, lag3: 0.2068, lag4: 0.3188, lag5: 0.3319, lag6: 0.3316, lag7: 0.0901
- **market\_daily\_pnl** → **Sentiment**:
  - lag1: 0.00213, lag2: 0.00383, lag3: 0.00368, lag4: 0.00822, lag5: 0.01480, lag6: 0.02242, lag7: 0.01225

### 8.3 Interpretation

- There is **stronger evidence that market moves drive sentiment** (`market_daily_pnl` → `sentiment`) than the reverse. P-values for this direction are consistently < 0.05.
- Sentiment does not robustly Granger-cause `market_daily_pnl` at a daily frequency ( $p > 0.05$  for most lags).
- **Implication:** The Fear/Greed index in this dataset appears to be a **reactive indicator**, following market activity rather than leading it.

## 9 Notebooks & Code Structure

- `1_eda.ipynb`: Loads cleaned data for basic exploratory data analysis, distributions, and plots.
- `2_features_and_aggregation.ipynb`: Demonstrates feature engineering (e.g., rolling features) and creates daily aggregate files.
- `3_analysis.ipynb`: Contains example code for OLS regression, the event study, and stubs for further analysis (Granger, VAR, ML).

## 10 Final Detailed Report Contents (`final_report_detailed.pdf`)

1. **Title & Executive Summary:** Key numeric findings and recommendations.
2. **Data & Cleaning Log:** Detailed list of all transformations with justifications.
3. **Exploratory Data Analysis:** Time-series plots, histograms, logarithmic tail visualizations, and full numeric stats.

4. **Rolling Correlation:** A 30-day rolling correlation plot and its interpretation.
5. **Event Study:** Numerical results and embedded charts for both market-level and account-level analyses with confidence intervals.
6. **Granger Causality:** Detailed ADF summaries, differencing decisions, p-value tables, and interpretation.
7. **Recommendations & Next Steps:** Prioritized list including panel regressions, clustering, and predictive modeling.
8. **Appendix:** Full file list and reproduction steps.

## 11 Limitations, Biases, and Validity Threats

- **Skew / Outliers:** Aggregate metrics are dominated by a few large PnL trades. Robust statistics are necessary.
- **Confounding (Omitted Variables):** The analysis lacks controls for BTC returns, volatility, or macro news, which could create omitted-variable bias.
- **Reverse Causality:** Granger tests indicate market moves precede sentiment changes, making naive predictive use of the index risky.
- **Sampling & Survivorship Bias:** Data is limited to Hyperliquid and may not represent the broader market.
- **Time Resolution Mismatch:** Daily sentiment may obscure intraday patterns where it could be predictive.
- **Future-Dated Timestamps:** The presence of future timestamps must be investigated and resolved.
- **Stationarity & Small Samples:** Account-level series were often too short for robust time-series analysis after differencing.
- **Metric Selection:** Summing PnL across traders (pooling) ignores heterogeneity. Panel models with fixed effects would provide better inference.

## 12 Recommended Next Analyses

### 12.1 Panel Fixed-Effects Regression

- **Purpose:** Control for account-specific traits and test if lagged sentiment predicts daily PnL after controlling for activity.
- **Specification:**

$$\text{daily\_pnl}_{it} = \alpha_i + \beta_1 \cdot \text{sent\_index\_norm}_t + \beta_2 \cdot \text{sent\_index\_norm}_{t-1} + \gamma \cdot X_{it} + \delta_t + \epsilon_{it}$$

Where  $\alpha_i$  is the account fixed effect,  $\delta_t$  is an optional day fixed effect, and  $X_{it}$  includes controls like `num_trades` and `BTC_return_t`.

- **Code Snippet (linearmodels):**



```

1 from linearmodels.panel import PanelOLS
2 # ... (load and merge data with sentiment and BTC prices)
3 df = df.set_index(['account', 'date'])
4 df['sent_lag1'] = df.groupby('account')['sent_index_norm'].shift(1)
5 exog_vars = ['sent_index_norm', 'sent_lag1', 'num_trades',
6             'avg_size', 'btc_return', 'realized_vol']
7 exog = df[exog_vars].dropna()
8 mod = PanelOLS(df['daily_pnl'], exog, entity_effects=True,
9               time_effects=False)
10 res = mod.fit(cov_type='clustered', cluster_entity=True)
11 print(res.summary)

```

## 12.2 Predictive Model (Classification)

- **Target:** profitable\_day (where daily\_pnl > 0).
- **Approach:** Use features like lagged sentiment, rolling stats, and BTC data. Train models (Logistic Regression, XGBoost) using time-aware cross-validation (TimeSeriesSplit) and evaluate with ROC-AUC.

## 12.3 Trader Clustering (Segmentation)

- **Features:** Use account-level stats (mean PnL, volatility, trade frequency) to cluster traders using KMeans. Analyze each cluster's sensitivity to sentiment separately.

## 12.4 Event Study Variants

- **Robustness:** Re-run studies with different thresholds (e.g., 5%, 20%), windows (e.g., [-1, +1]), and use robust metrics like medians.

## 13 How to Incorporate BTC Price & Volatility

1. **Get Data:** Obtain daily BTC closing prices aligned with the sentiment data dates.
2. **Compute Returns:** Calculate daily returns:  $\text{btc\_return}_t = (\frac{\text{close}_t}{\text{close}_{t-1}} - 1)$ .
3. **Compute Volatility:** Calculate realized volatility, e.g., the rolling 30-day standard deviation of daily returns.
4. **Merge:** Merge these new features into the daily aggregate datasets on the date column.

## 14 Full Prioritized Action Plan

- **Priority 0 — Data Integrity:** Confirm and fix or remove future-dated timestamps.
- **Priority 1 (2–4 hours):**
  1. Run the panel fixed-effects regression with BTC controls.
  2. Conduct robust event studies using medians.
- **Priority 2 (1 day):**

1. Build the classification model for predicting a profitable day.
2. Cluster traders and analyze sentiment sensitivity per cluster.

- **Priority 3 (2–3 days):**

1. Develop a full Streamlit dashboard with interactive filters and visualizations.
2. Prepare the final slide deck with the top 3 insights and actionable recommendations.

## 15 Concrete Recommendations for a Trading Team

1. **Do not use the raw Fear/Greed index as a leading signal.** Evidence suggests it is a reactive indicator. Use it conditionally, after controlling for price and volatility.
2. **Use robust risk-sizing during periods of Extreme Greed.** Event studies show elevated PnL, which likely corresponds to higher volatility. Consider reducing position sizes or tightening stop-losses.
3. **Segment traders.** Identify clusters that are particularly sensitive to sentiment changes and manage their risk accordingly.
4. **Use robust performance metrics.** Monitor daily performance using medians or winsorized means to avoid being misled by outliers.
5. **A/B test any new policy.** Before deploying a sentiment-aware strategy, backtest it rigorously using time-series cross-validation.

## 16 Exact Code Snippets & Helper Utilities

### 16.1 Auto-detect DATA\_ROOT

```

1 from pathlib import Path
2 BASE = Path.cwd()
3 if (BASE / 'submission_package').exists() and not (BASE / 'trades_clean.csv').exists():
4     DATA_ROOT = BASE / 'submission_package'
5 elif (BASE / 'trades_clean.csv').exists():
6     DATA_ROOT = BASE
7 else:
8     DATA_ROOT = BASE / 'submission_package'
9 print("DATA_ROOT:", DATA_ROOT)

```

### 16.2 Load Cleaned Data

```

1 import pandas as pd
2 trades = pd.read_csv(DATA_ROOT / 'trades_clean.csv', parse_dates=['time'])
3 sent = pd.read_csv(DATA_ROOT / 'sentiment_clean.csv', parse_dates=['date'])
4 daily = pd.read_csv(DATA_ROOT / 'daily_account_agg.csv', parse_dates=['date'])
5 merged = pd.read_csv(DATA_ROOT / 'merged_daily_sentiment.csv', parse_dates=['date'])

```

## 16.3 Event Study Helper

```
1 import numpy as np
2 def build_event_matrix(event_dates, series_df, value_col, window=3):
3     rows=[]
4     series_map = series_df.set_index('date')[value_col].to_dict()
5     for ev in event_dates:
6         for k in range(-window, window+1):
7             d = (pd.to_datetime(ev) + pd.Timedelta(days=k)).date()
8             rows.append({'event_date': ev, 'rel_day': k,
9                         'value': series_map.get(d, np.nan)})
10    return pd.DataFrame(rows)
```

## 16.4 Granger Causality Quick Call

```
1 from statsmodels.tsa.stattools import grangercausalitytests
2 data = merged[['market_daily_pnl', 'sent_index_norm']].dropna()
3 # ensure stationarity or diff once
4 grangercausalitytests(data[['market_daily_pnl', 'sent_index_norm']], maxlag=7)
```

## 17 Location of Important Outputs

- **Detailed Report:** /mnt/data/submission\_package/final\_report\_detailed.pdf
- **Cleaned Data:** /mnt/data/submission\_package/\*.csv
- **Event Study Plots:** /mnt/data/submission\_package/event\_study\_\*.png
- **Granger Results:** /mnt/data/submission\_package/granger\_full\_summary.json
- **Notebooks:** /mnt/data/submission\_package/\*.ipynb